Machine learning in battery management systems: a two-tier scoping review of 799 eligible studies with in-depth classification and validation audit

Hussain Touhid Siddiquee, Syeda Salsabil Islam Ariya, and Jasimul Islam Chowdhury

***Abstract*—Machine learning (ML) has moved from a promising addition to battery research to a load-bearing component of battery management systems (BMS), but the validation practices behind the accuracy claims are rarely examined. A structured scoping review with a two-tier evidence base is presented: an *eligible corpus* of 799 peer-reviewed journal articles (2019–2026), selected by uniform rule-based criteria, and an in-depth *classified subset* of 103 studies coded by function, method family, chemistry, data source, and validation practice. Union-share bibliometrics across the eligible corpus locate the center of gravity of the screened literature in health estimation (SOH, 20% of records), state estimation (SOC, 17%), and prognostics (13%), with charging, thermal management, balancing, and vehicle- or fleet-level work relatively under-served for a transportation-electrification readership. Within the classified subset, three weaknesses recur. Only 8 of 88 primary-application papers evaluate online or on real vehicles; 17 of 88 rely on within-cell or k-fold splits that leak trajectory information; and the largest single data source is private laboratory cells (36 of 88), a reproducibility constraint at least as limiting as the much-cited benchmark concentration (23 papers on three datasets). Validation practice does not demonstrably differ between physics-informed/hybrid and purely data-driven papers under either definition (36% versus 52% protocol-only), and a label-noise sensitivity analysis brackets the strong-protocol share between 35% and 60%. The paper closes with a requirement-level mapping of BMS functions onto UN GTR 20, GB 38031, and ISO 26262, a positioning against prior reviews, and a twelve-item reporting checklist; all data and code are released.**

***Index Terms*—machine learning, battery management systems, lithium-ion batteries, state of charge, state of health, remaining useful life, deep learning, physics-informed neural networks, transfer learning, battery safety**

I. Introduction

A. The battery management problem

Lithium-ion batteries now anchor three converging transitions — electrified road transport, renewable-heavy stationary storage, and the electrification of aerospace and marine systems — and in each the battery is an actively managed asset whose usable capacity, lifetime, and safety depend on real-time operating decisions. The battery management system (BMS) is the embedded cyber-physical layer that makes those decisions — sensing, state estimation, protection, balancing, and charging negotiation [1]. The difficulty stems from the fact that the quantities that matter most — state of charge (SOC), state of health (SOH), internal temperature, and the onset of degradation or fault mechanisms — are not directly measurable and must be inferred from noisy, partially observable terminal measurements under widely varying conditions. Model-based approaches built on equivalent-circuit or electrochemical models face a persistent fidelity-versus-tractability trade-off, and their parameters drift as cells age [2]. It is in this inference-and-decision layer that machine learning has made its most forceful entry, and the stakes are highest in transportation electrification: electric-vehicle packs face automotive safety expectations — UN GTR 20 on board energy storage rechargeable systems, GB 38031 (GB/T 38031) on traction battery safety requirements, and ISO 26262 on functional safety of electric/electronic systems — that make state estimation, degradation prognosis, and fault detection safety-relevant BMS functions, not research add-ons. Section IX returns to these three standards and maps each estimation function to the property the standard actually demands (detection reliability, gradient/temperature limits, protection timing) — the mapping that turns the validation findings of this review into vehicle-level stakes.

Machine learning entered battery management through estimation on hand-engineered features. Two studies then reset expectations. Severson and colleagues predicted the entire cycle life of lithium-ion cells with a median test error near 9% from the first hundred cycles alone, on a public dataset of 124 cells [3]. Attia and colleagues demonstrated closed-loop Bayesian optimization of fast-charging protocols, identifying high-cycle-life protocols among 224 candidates in 16 days of cycling rather than the more than 500 days an exhaustive search would take [4]. The five years that followed produced applications across every BMS function: recurrent and convolutional networks for SOC and SOH [5] [6], graph and transformer architectures for prognostics [7] [8], physics-informed neural networks embedding electrochemical and thermal priors [9] [10], transfer learning across chemistries and vehicles [11] [12], reinforcement learning agents [13], and, most recently, LLM-based studies of station-side incident-risk analysis and vehicle-cloud management [14] [15].

Manuscript received September 2026. (Corresponding author: Hussain Touhid Siddiquee.)

The authors are with Leading University, Sylhet, Bangladesh. The complete evidentiary record of this review—classified corpus, screening provenance log, protocol audit, codebook, and the blind inter-coder reliability study, together with the full analysis pipeline—is released as a Zenodo deposit, DOI: 10.5281/zenodo.23436202; see Data and Code Availability. Item 12 of the reporting checklist this paper proposes makes such a DOI-backed deposit the reporting norm.

B. Existing reviews and the literature gap

The review literature has grown alongside the applications but is fragmented along the same axes as the primary studies: function-specific reviews cover SOC estimation [16] [17], SOH estimation [18], and prognostics [19]; method-specific reviews follow deep-learning state estimation [20], physics-informed impedance diagnostics [21], or interpretable prognosis [22]; and domain reviews examine fault diagnosis [23] [24] and thermal management [25], alongside broader degradation-method surveys [26] and EV-powertrain ML reviews with BMS scope [27]. What is missing is a cross-functional account of the field's trajectory: which method families are used for which tasks, on which data, validated how — and, among existing cross-functional surveys, none quantifies validation practice at all. Table VII in Section IX positions this review against its closest competitors and states the added value explicitly; the short version is that prior reviews synthesize *what is claimed*, while this review additionally measures — within a 103-study classified subset — how those claims are validated, ties both to three safety standards, and releases the record-level dataset to make every count auditable.

C. Contribution of this paper

This paper makes four contributions: (i) a rule-based *eligible corpus* of 799 peer-reviewed journal articles (2019–2026), selected by uniform criteria applied to all screened Crossref records — no function quotas, year caps, or citation thresholds — so that field-level shares are measured, not curated, with the retrieval, screening, and tagging provenance of every record released; (ii) an in-depth *classified subset* of 103 studies from that corpus, coded along eight axes (function, method family, chemistry, data source, validation practice, and more) with the classification dataset, codebook, per-paper protocol audit, and a 30-paper blind inter-coder reliability study released together with a label-noise sensitivity analysis of the headline counts; (iii) a two-tier synthesis — measured bibliometrics from the eligible corpus (Section III) plus function-by-function critical review of the classified subset (Sections IV–VIII) — with a BMS-function-to-safety-standard mapping and a gap-driven research agenda including a twelve-item reporting checklist (Section IX); and (iv) a transportation-electrification cut of the corpus (Section III-B) separating cell-level laboratory evidence from pack-, vehicle-, and fleet-level practice. Fig. 1 previews the classified subset, mapping it across management functions and method families: the center of gravity sits in the health-prognostics complex, hybrid model–data approaches span every function, and learning-based charging control remains a concentrated Bayesian-optimization stronghold. The remainder of the paper is organized as follows. Section II details the review methodology and its limitations. Section III presents the two-tier corpus overview, five-year bibliometrics, and the transportation-electrification cut. Sections IV–VIII review machine learning for state estimation, health estimation and prognostics, charging control, thermal management, and fault detection and safety. Section VIII examines cross-cutting directions, and Sections IX and X give the research agenda and conclusions.

II. Review Methodology

A. Search strategy

All candidate records were retrieved programmatically from the Crossref REST API between 4 and 5 September 2026 using 30 bibliographic query strings spanning the application areas above. Queries were issued in five passes: a citation-sorted 2020–2026 pass, dedicated 2020–2022 and 2023 windows, a mid-window 2021–2023 coverage pass, and a relevance-sorted recency pass starting June 2024, yielding 3,000 records; after DOI-based deduplication (1,450 duplicates removed), 1,550 unique records remained. Six landmark studies — including [3] and [2] — were additionally safeguarded as forced seeds: five were retrieved and verified individually because citation-sorted retrieval had crowded them out, and one was already among the top-cited hits. Three properties should be stated plainly. First, the Crossref registry is a metadata registry, not a discovery engine like Scopus or Web of Science; it was chosen so that every included record is individually verifiable by direct DOI resolution, at the cost of a recall floor. Second, citation-ordered passes favor established English-language work; all trend statements therefore describe this citation-weighted, journal-article sample rather than the field in the aggregate. Third, January–May 2024 is a known retrieval hole — the citation-sorted pass returns almost nothing from those months and the recency pass starts at June 2024 — so the 2024 dip visible in year distributions is a design artifact of this retrieval, not evidence about publication volume; per-year shares in Section III-A are therefore reported as shares of Crossref-reported query totals (the full result count of the same queries per year), not of the capped retrieval.

B. Screening and rule-based selection

The screening funnel is shown in Fig. 2. A title-level relevance gate requiring both battery-related and machine-learning terminology retained 883 records and excluded 667 (predominantly adjacent-field papers that surfaced through loose matching); the gate is a mechanical keyword rule, and its term lists are released (tagging_terms.json) together with all 1,550 screening decisions (screening_log.csv). Venue and document-type screening removed 84 records — conference abstracts and proceedings series, preprint servers, and non-scholarly venues — leaving 799 eligible records. Selection then proceeds in two tiers, replacing the single-quota design that an earlier draft of this review used. *Tier 1 (eligible corpus).* Every one of the 799 eligible records is included, subject only to venue screening; no function quotas, year caps, or citation thresholds are applied. Every one of the 799 eligible records carries a Crossref-registered DOI, and every record carried forward to Tier 2 full-text coding was individually DOI-resolved. Function tagging at this tier is a rule-based keyword pass over titles (term lists released); 202 of 799 records (25%) carry no management-function tag — most are method-level or cross-cutting records (physics-informed, second life, digital twins, benchmark papers) — so per-function shares below are computed over tagged records and cross-function totals are reported as union shares, never as sums

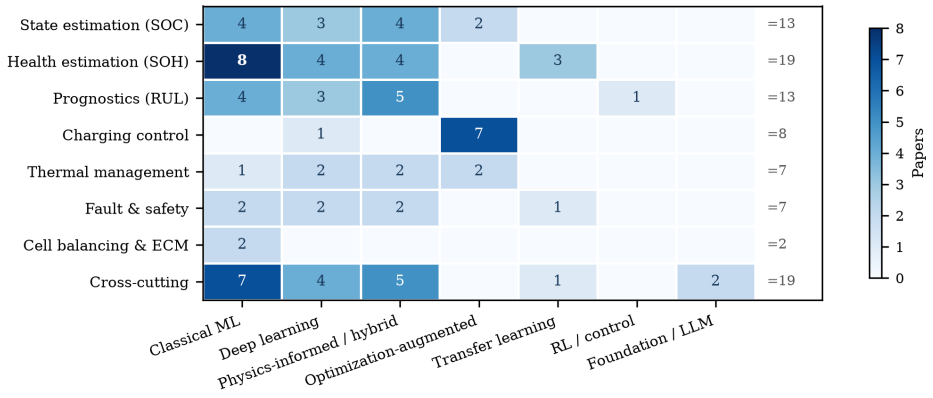


Fig. 1. The classified subset mapped across battery-management functions and method families (the 88 primary-application papers). The remaining 15 classified papers are 12 method-coded reviews ("Review / perspective") plus 3 review-type syntheses coded under their dominant method family — two of which sit in the physics-informed column of Table I and are labeled in Table II. Rates that use the primary denominator are 88, not 103.

of per-tag percentages. *Tier 2 (classified subset).* Because in-depth coding of chemistry, data source, and validation practice requires full-text assessment, full-text coding was applied to a 103-study subset of the eligible corpus, curated with the same eight-function stratification used by the tagging pass, with the 2019 field-defining seed (Severson et al.) and six landmark studies ([4], [28], and four others) included as forced entries. This subset is curated, not sampled: the released record-level log gives the rule or rationale for every study included or dropped, and its function and year composition does not match the eligible corpus's measured shares exactly. Section III therefore never generalizes subset distributions to the field; where the subset is the only source (method family, validation practice), the text reports subset statistics and their selection caveats explicitly.

C. Classification framework, validation definitions, and limitations

Each paper was classified along six fields: management function (eight values, including a cross-cutting cluster); dominant method family (classical ML, deep learning, physics-informed/hybrid, optimization-augmented, reinforcement/control learning, transfer learning, foundation models/LLMs, review); chemistry; primary data source; validation approach; and a coded summary of the headline finding and principal limitation. Three terms are used consistently: an *eligible record* is one of the 799 Tier-1 studies; a *classifiable paper* is one of the 103 Tier-2 studies; and a *primary-application paper* is one of the 88 classified papers whose contribution is a primary application — the 12 method-coded reviews and 3 review-type syntheses coded under their dominant method family are excluded from every rate that uses this denominator (Table II labels each such paper explicitly). Validation practice is coded against four definitions used throughout this paper: (1) *within-cell split* — train and test segments drawn from the same cycling trajectory; (2) *unsupervised k-fold cross-validation* — folds assigned over samples of the same trajectories without cell isolation; (3) *cross-cell holdout* — every test cell held out entirely from training (with cross-chemistry and cross-domain holdout extending this isolation across chemistry or application); and (4) *online or real-vehicle/fleet validation* — the deployed system evaluated on live vehicle data. Cross-cell, cross-chemistry, cross-domain, and online/fleet are together the *strong protocols*; the four remaining coded environments (experimental hardware, simulation plus hardware spot-check, simulation-only, case study) characterize where a study was exercised but say nothing about train–test isolation, and they are reported separately — never mixed into strong-protocol percentages. Borderline assignments, the absorbing rule for ambiguous cases (a study is credited with the *strongest* protocol it reports for its headline claim), and the decision tree with examples are given in the released codebook. Classification was performed by the lead author with conservative coding from full texts; to probe the reliability of that coding, two independent coders later blind-re-coded a 30-paper stratified random sample (drawn by function and year band, seed-pinned and reproducible) under the released codebook, without access to the released labels, to the population statistics, or to each other. Blind agreement between the two independent coders is strong on three of the five classification fields: management function (κ = 0.92, Krippendorff α = 0.92), method family (κ = 0.88, α = 0.88), and primary data source (κ = 0.88, α = 0.88). Two fields fell below the $\kappa \geq 0.66$ threshold: chemistry (κ = 0.55) and validation protocol (κ = 0.48). Part of the shortfall is an evidence-access artifact: 23 and 28 of the 30 sample papers (respectively) could be coded from abstracts alone, and these two fields are precisely those least often stated at abstract level. But it must be reported straight that the aggregation Section III's counts rest on does not escape this: on strong-protocol versus non-strong, the two coders agree on 25 of 30 papers (raw agreement 83%, κ = 0.52 at the sample's prevalence — below the 0.66 threshold the paper itself applies). The released labels are therefore treated as one plausible coding, and all headline validation counts in Section III carry a scenario sensitivity analysis built from the observed coder-versus-label disagreement in this sample: the central strong-protocol estimate (48% of the 88 primary

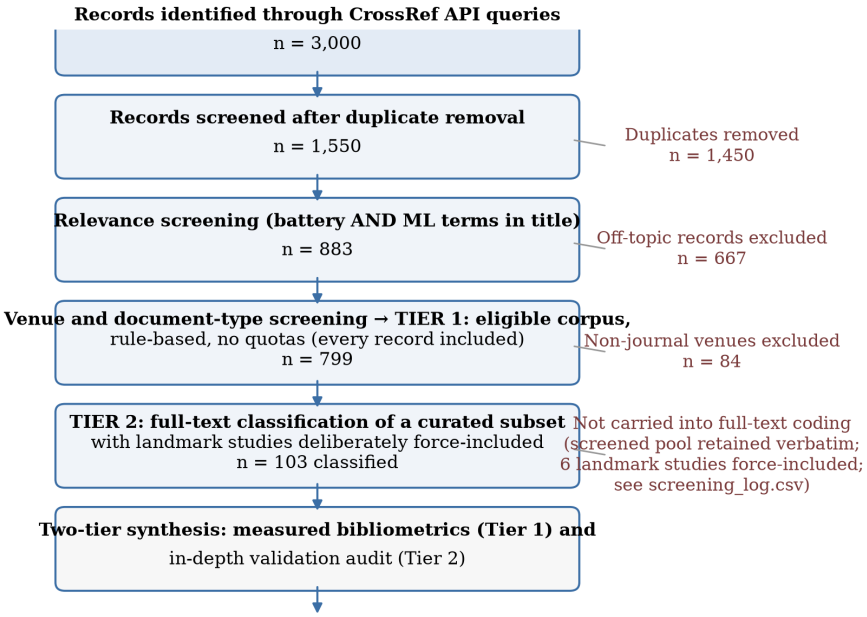


Fig. 2. Study-selection flowchart with the two-tier design. Records were retrieved from the Crossref REST API on 4–5 September 2026 (five passes; January–May 2024 retrieval hole documented in Section II-A), deduplicated, screened at title level by a released mechanical keyword rule (883 retained, 667 excluded), then at venue and document-type level. Tier 1 includes every eligible record (799, rule-based, no quotas); Tier 2 carries a curated full-text-classified subset of 103 with six landmark studies included deliberately and five force-included individually from outside the pool. Record-level provenance is released as screening_log.csv (all 1,550 decisions) and tagging_terms.json.

papers) spans 35–60% under plausible label error, and no qualitative conclusion below depends on where in that range the truth lies (Section III-C). Coder access constraints, training, and the complete per-paper, per-field disagreement table are released together with both locked coding sheets and the hash-tracked log; protocol-versioning candidates (a "protocol not reported" code; a chemistry inference rule) are logged as candidates and alter no released label. The synthesis proceeds on two tracks: measured bibliometrics from the eligible corpus and cross-tabulations from the classified subset, both generated programmatically from the released datasets (Section III), and qualitative function-by-function critique (Sections IV–VIII). The methodology's limitations are acknowledged: Crossref coverage excludes influential conference papers; title-level screening can miss relevant work; citation-ordered retrieval privileges established English-language venues; the per-year retrieval window has one documented hole (Section II-A); the Tier-2 subset is curated rather than sampled, so validation-practice rates are subset statistics that do not automatically generalize to the eligible corpus or the field; and reliability is documented on a 30-paper blind re-coding subset rather than measured corpus-wide (three of five fields above threshold, two below, with the resulting sensitivity ranges reported). DOI resolution, moreover, verifies that a record exists, not that its venue is topically sensible: one corpus study (a physics-informed SOC network, [10]) appears in an RFID-special-issue journal — we flag it here rather than dropping it, and the released audit file (final_audit.json) records every reference's Crossref-resolved title and venue so that misattributions of this kind are checkable rather than hidden. These are the reasons this paper describes itself as a scoping review with bibliometric classification rather than a meta-analysis.

III. Corpus Overview and Five-Year Evolution

A. Publication trend and method-family evolution

Two layers of trend evidence are reported. On the *measured* tier (eligible corpus, 799 records), the capped retrieval captures 55, 85, 139, and 223 journal articles for 2020–2023, a documented retrieval hole totaling 56 in 2024 (Section II-A), then 126 and 114 in 2025 and — partial year, retrieval date September 2026 — 114 in 2026 (Fig. 3, left panel). These are counts within a capped retrieval, not census totals, and they do not support absolute-growth claims ("the literature quadrupled" would require an uncapped retrieval). A per-year external denominator was tested and discarded: Crossref's reported result totals for the same 30 queries (a published script and its output ship with the deposit) rely on a loose token-level bibliographic match whose totals run into the hundreds of thousands and vary strongly with query phrasing, so they are not interpretable as field denominators. What remains defensible is direction within the sample: the eligible-corpus counts rise through 2020–2023, dip at the retrieval hole, and remain high in 2025–2026, and the tagged composition of those records shifts toward estimation and hybrid methods (next paragraph).

On the *classified subset*, the defining finding is the compositional shift in method families (Fig. 3, right panel; Table I). Classical ML — Gaussian process regression, support vector machines, ensembles, shallow networks on engineered features — contributes 11 of the 23 subset papers published from 2019 through 2021, while physics-informed and hybrid model–data learning rises year by year (1, 2, 3, 3, 4, 4, and 5 primary papers in 2020–2026) to become the modal family by 2024–2026 — 11 of its 22 primary papers fall in 2024–2026, and it is the largest component of the 7 papers published in 2026 (Table II gives corpus-wide totals). Optimization-augmented learning (Bayesian optimization, surrogate-assisted and metaheuristic schemes) peaks in 2023–2024 (11 papers); transfer learning accelerates in 2024–2025 (5 papers); reinforcement learning appears as the dominant family in 1 paper and as an embedded component in at least one more [13] [29]; the first foundation-model and LLM-based studies appear in 2025–2026 (2 papers) [14] [15], with preprint and conference LLM-based battery estimation beginning earlier outside the screened venues. These family shares describe the subset only; the subset is curated, and small-denominator years carry wide sampling error. Two qualifications matter: the families are not mutually exclusive (a PSO-tuned LSTM is both deep learning and optimization-augmented), so the composition understates method mixing; and physics-informed papers are, on average, validated no more stringently than others even under the most permissive definition of stringency — a point quantified, with its sensitivity ranges, in Section III-C.

B. Measured application areas, the transportation-electrification cut, and venues

Because function tags were applied mechanically to all 799 eligible records, the application-area distribution is *measured* rather than quoted from a curated sample. Tagging is a rule-based keyword pass over titles with released term lists; 597 records (75%) carry at least one management-function tag, while the remaining 202 are method-, second-life-, digital-twin-, or benchmark-level records without a specific BMS function. Of all eligible records, health estimation (SOH) leads at 20% (158 records; 27% of those function-tagged), followed by state estimation (SOC) at 17% (133), prognostics (RUL) at 13% (103), thermal management at 12% (99), fault and safety at 12% (96), charging control at 8% (64), and cell balancing at 5% (38). These per-tag percentages must not be summed — a record can carry several tags; the union of SOC, SOH, and RUL covers 46% of eligible records, and adding the four control-and-safety functions brings management-function coverage to 75%. The estimation concentration is a measured pattern consistent with the commercial pull of lifetime prediction for warranties and second-life markets, and is mirrored (with a heavier cross-cutting share) in the classified subset (Fig. 1; Table II gives subset composition by year; Table IV maps representative studies). Two measured concentrations matter specifically for transportation electrification. First, the electrification cut: only ≈2% of eligible records ($n = 17$) are tagged fleet- or cloud-orchestration-adjacent, and in the classified subset only 10 of 88 primary papers use real vehicle or fleet data as their primary source against 36 resting on laboratory cells — the classified literature is predominantly cell-level and laboratory-scale, one integration level below the pack- and vehicle-level questions that UN GTR 20, GB 38031, and ISO 26262 pose. Second, venue concentration is marked: in the eligible corpus the Journal of Energy Storage (103 records), Energy (61), and Batteries (45) lead, and the same two energy-storage journals head the classified subset (13 papers each, with Applied Energy third at 8) — energy-storage venues, not transportation-electrification venues, absorb the field's output. The thin representation of cell balancing is also telling: balancing remains dominated by classical control practice, and the subset's most-cited balancing-related result is the demonstration that routine balancing can be dispensed with in well-managed packs, reducing hazardous lithium-plating side reactions [30].

C. Data landscape and validation practice

The data underpinning the classified subset are strikingly concentrated: 36 of the 88 primary-application papers rest on laboratory cycling of the authors' own cells; the public Severson/Toyota–MIT–Stanford (MATR) dataset alone supports 11, NASA and CALCE 6 each, and laboratory impedance spectroscopy 6; 12 are simulation-based. Only 10 papers use real vehicle or fleet data as their primary source, and *data use* is a distinct measure from *validation protocol*: 8 of the 10 fleet-data papers are validated online or on real vehicles, and the other two through cross-domain holdout or case-study protocols (Table III). Validation practice is arguably the classified subset's most consequential weakness, and it is now defined and audited explicitly. Per the definitions of Section II-C, cross-cell holdout — the strongest common laboratory protocol — accounts for 24 primary-application papers, cross-chemistry holdout for 7, cross-domain holdout for 3, and online or fleet validation for 8; together these *strong protocols* cover 42 of the 88 primary-application papers (48%). Within-cell splits (12 papers) and unsupervised k-fold cross-validation (5) leak trajectory information between train and test, producing optimistic errors that do not survive cell-to-cell variability; the remaining 29 studies establish hardware operation (20), simulation- or case-level evidence without a generalization protocol (9) — 42+17+29 = 88. The complete per-paper protocol assignment is released in papers.csv (valid field) and tabulated as a supplement (protocol_audit.csv), so the 42-of-88 arithmetic above can be verified row by row, and 30 of the 88 assignments were re-derived blind by two independent coders (Section II-C).

Crossing validation practice with method family, the cautious reading is: 8 of the 22 physics-informed/hybrid primary papers (36%) use a strong protocol, against 34 of the 66 other primary papers (52%) — Fisher's exact $p = 0.32$, odds ratio 0.55 (95% CI 0.21–1.46). Under the earlier, broader definition that also counted experimental-hardware demonstration as strong validation, the contrast is 13 of 22 (59%) versus 49 of 66 (74%), OR 0.50 (95% CI 0.19–1.36, $p = 0.19$). We report both and claim neither as a model-family effect: the difference is not statistically significant under either

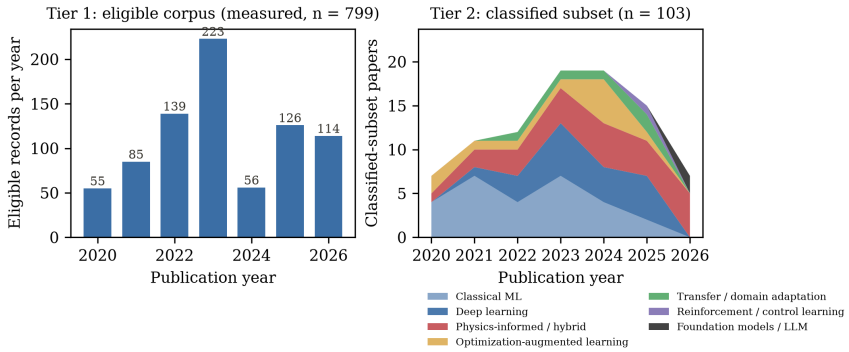


Fig. 3. Two-tier trend evidence. Left: Tier 1 eligible-corpus records per year across the capped retrieval (the 2024 dip is the documented January–May 2024 retrieval hole, not a decline in the field). Right: Tier 2 classified-subset method-family papers by year (curated subset; wide sampling error in small years; the 2019 seed is omitted for scale).

TABLE I

METHOD FAMILIES: REPRESENTATIVE TECHNIQUES, STRENGTHS, AND LIMITATIONS. COUNTS ARE THE 88 PRIMARY-APPLICATION PAPERS; THE THREE REVIEW-TYPE SYNTHESES CODED UNDER A METHOD FAMILY PLUS THE 12 METHOD-FAMILY REVIEWS ARE ACCOUNTED FOR IN TABLE II.

Method family	**Papers**	**Representative techniques**	**Strengths**	**Limitations**
Classical ML	28	Gaussian process regression, SVM/SVR, ensembles, shallow networks on engineered features	Calibrated uncertainty; small-data friendly; cheap to embed	Feature engineering is manual; ceilings under complex dynamics
Deep learning	19	LSTM/GRU, CNN, attention/Transformer, autoencoders, graph networks	Learns features from raw curves; strong trajectory modeling	Data-hungry; weak extrapolation; opaque
Physics-informed / hybrid	22	PINNs, physics-constrained autoencoders, filter-embedded networks, multiphysics fusion	Structure aids extrapolation; reduced data demand	Inherits simplified-model bias; harder to train
Optimization-augmented	11	Bayesian optimization, surrogate-assisted multi-objective design, meta-heuristic-tuned networks	Sample-efficient protocol discovery; explicit trade-off handling	Offline tuning cost; surrogate bias
Transfer / domain adaptation	5	Domain-adversarial training, fine-tuning across chemistry/protocol	Directly targets the generalization gap	Negative transfer; needs shared protocols
Reinforcement / control learning	1	Deep-RL prognostic model selection (also an embedded component in further studies)	Adaptive decision-making along trajectories	Training stability; scarce hardware validation
Foundation models / LLM	2	Domain-tuned LLM agents for risk analysis and orchestration	Reasoning over heterogeneous text/system context	Hallucination risk; latency and cost
Review / perspective	12	Narrative and systematic syntheses	Maps the field; identifies shared gaps	No primary validation

TABLE II

CORPUS COMPOSITION: MANAGEMENT FUNCTION BY PUBLICATION YEAR. THE REVIEWS COLUMN COUNTS THE 12 METHOD-FAMILY REVIEWS PLUS THE 3 REVIEW-TYPE SYNTHESES CODED UNDER THEIR DOMINANT METHOD FAMILY; PRIMARY = TOTAL − REVIEWS IS THE 88-PAPER PRIMARY-APPLICATION COUNT USED IN FIG. 1.

Function	**'19**	**'20**	**'21**	**'22**	**'23**	**'24**	**'25**	**'26**	**Total**	**Reviews**	**Primary**
Cell balancing & ECM	0	0	0	0	1	1	0	0	2	0	2
Charging control	0	2	0	0	2	3	1	0	8	0	8
Cross-cutting	0	1	0	6	8	3	2	4	24	5	19
Fault & safety	0	0	2	0	2	3	2	0	9	2	7
Health estimation (SOH)	0	2	3	2	4	3	4	2	20	1	19
Prognostics (RUL)	1	1	3	2	1	1	6	0	15	2	13
State estimation (SOC)	0	3	4	2	1	5	2	0	17	4	13
Thermal management	0	0	1	1	3	1	1	1	8	1	7
Total	1	9	13	13	22	20	18	7	103	15	88

definition, confounding is unadjusted (physics-informed work skews toward early-stage and simulation-heavy studies), and n = 22 is small. A label-noise sensitivity analysis built from the observed blind coder-versus-label disagreement (Section II-C) moves the overall strong-protocol share across 28–61 of the 88 papers and leaves the method-family contrast non-significant in 95% of scenario draws, so the corpus's qualitative conclusions survive plausible coding error quantitatively. The honest reading remains: a physics prior does not *demonstrably* buy validation rigor — and settling the question requires the standardized protocols proposed in Section IX, not more corpus cross-tabs.

The data concentration can be stated as a two-pole fragmentation metric (Fig. 4). Across 2019–2026, 23 primary papers — 26% of the corpus — rest on the same three public benchmarks (Severson/MATR, NASA, CALCE), whose combined cell count is on the order of a few hundred; meanwhile, the single largest source overall is private laboratory cycling (36

TABLE III
CROSS-TABULATION OF FLEET-DATA USE AGAINST ONLINE VALIDATION (103 CLASSIFIED PAPERS). FLEET-DATA USE DESCRIBES THE PRIMARY DATA SOURCE; ONLINE VALIDATION DESCRIBES THE VALIDATION PROTOCOL; THEIR OVERLAP IS 8 OF 10 FLEET-DATA PAPERS. A LABORATORY-DATA ONLINE DEMONSTRATION WAS RECODED TO EXPERIMENTAL HARDWARE AFTER FULL-TEXT RE-CHECK (SECTION II-C). THE RIGHT-HAND COLUMN POOLS THE ENVIRONMENT-ONLY VALIDATION CATEGORIES (SECTION III-C).

	Online / real-vehicle validation	**All other validation protocols (Table 5)**	**Total papers**
Real-vehicle / fleet data	8	2 (cross-domain holdout; case study)	10
Other data sources	0	93	93
Total	8	95	103

papers) whose data are rarely released. The empirical base is therefore not one benchmark monoculture but two poles: a small set of shared benchmarks onto which comparative claims crowd, and many unshared laboratory datasets over which accuracy claims cannot be audited at all. On the shared benchmarks, dataset-specific failure modes and implicit test-set familiarity inflate apparent comparability; on the private datasets, reproducibility is unavailable in principle. Both faces of this pattern — the empirical base on which the community calibrates and compares estimators remaining approximately fixed while the literature builds on top of it — we call benchmark saturation: the pattern this review flags as the most reproducibility-relevant feature of the classified subset's data landscape.

IV. MACHINE LEARNING FOR BATTERY STATE ESTIMATION

SOC estimation is the most mature BMS function, and the corpus shows how machine learning entered it. The early pattern pairs learning with classical estimation: Gaussian process regression for pack-level SOC with native uncertainty quantification [31], and — most influentially — recurrent networks embedded in unscented Kalman filters, where the LSTM supplies a learned observation model and the filter recursive correction [5]. A second generation tuned network structure itself: particle-swarm-optimized LSTM improved accuracy and noise robustness [32], and evolutionary-assisted CNN–BiLSTM hybrids report sub-1% drive-cycle errors [33]. Simulation-informed training pipelines substitute automotive-accurate simulated data for exhaustive cell testing [34], comparative benchmarks on standard drive cycles weigh ML regression families [35], pack heterogeneity is addressed by joint pack-SOC and cell-inconsistency estimation with balancing-aware formulations [36], and unified frameworks estimate SOH and SOC jointly from partial charging data [37], and large-format applications extend estimation to containerized storage strings [38], with comparative benchmarks weighing accuracy against embedded compute budgets [39] [40].

The most substantive recent shift is toward hybrid schemes in which a filter, an equivalent-circuit model, or a physics prior constrains the learner: GA-optimized LSTM coupled with improved adaptive Kalman filtering reaches low error under measurement noise that breaks pure-network estimators [41]; equivalent-circuit-informed ML on impedance features estimates SOC across temperatures [42]; and physics-informed recurrent structures embed battery dynamics directly into the training objective [10]. These schemes trade engineering complexity for extrapolation behavior — the property that matters most onboard. Yet 11 of the 13 subset SOC primary papers rest on the authors' own laboratory cells, one on laboratory impedance measurements, 9 of the 13 on a single cathode chemistry (NMC), and the one operation-scale demonstration — a stationary containerized storage system, not a vehicle — [38]; benchmark heterogeneity makes cross-paper comparison nearly meaningless [20] [17]. Accuracy on laboratory cells is no longer the binding constraint in this cluster; guaranteed accuracy on arbitrary packs over full cell aging, which the onboard estimator would need, is not demonstrated anywhere in the subset.

V. MACHINE LEARNING FOR HEALTH ESTIMATION AND PROGNOSTICS

SOH estimation and RUL prediction form the largest functional block of the corpus (32 primary papers: 19 SOH, 13 RUL). The founding pattern is classical: extract health-sensitive features from partial charge or discharge curves, then regress with Gaussian process or similar shallow models — incremental-capacity analysis with explicit uncertainty bounds [43], differential thermal voltammetry [44], and Gaussian-process classification of impedance spectra that identifies degradation modes [45], and deep Gaussian processes coupling prognosis with degradation-mode identification [46]. Measurement relaxations followed: random partial charging data proved sufficient for accurate SOH [47], and charging-segment adjustment refined features for vehicle-level use [48]. The landmark systematization is Roman et al.'s 179-cell pipeline showing that relaxed charge-curve features support accurate SOH estimation across a heterogeneous commercial fleet [28].

Deep architectures entered as sequence models over capacity-fade trajectories and representation learners over raw curves: graph convolutional networks fusing multiple health indicators [7], graph networks over partial discharge curves [49], BiLSTM–Transformer hybrids under fast-charging stress [8], hybrid-fusion networks for real-time prediction streams [50], and interpretable deep learning with transparency mechanisms [51]. Field-data studies absorb real-vehicle noise [6] and estimate pack-level SOH from fleet streams [52], while deep models estimate SOH without additional degradation experiments by learning trajectory features that generalize across chemistries [53]. Early-life prediction — forecasting lifetime from the first cycles — compresses test times from months to days [3] [54] [55], hybrid physics-plus-learned models give early RUL from limited cycles [56], extends to lithium-metal chemistries [57], and now runs on embedded boards in real time [58].

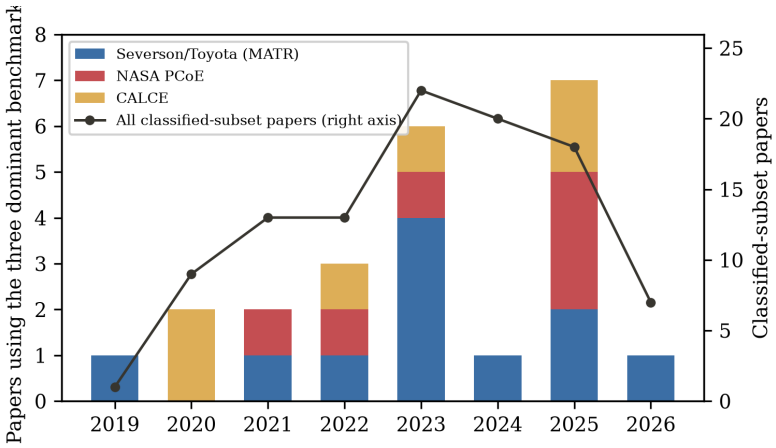


Fig. 4. Benchmark reuse versus classified-subset output. Papers in the classified subset using the three dominant public benchmarks per year (stacked bars) against classified-subset papers per year (line). Both quantities describe the curated 103-study subset, not the eligible corpus.

The prognostics cluster shows the clearest physics-informed turn: physics-informed online degradation models [59], hybrid reinforcement learning that reduces the data demand of SOH estimation [29], a generalizable PINN from limited early-cycle data [60], EIS-informed physics-constrained networks [61], a generic physics-informed RUL framework for small early-life samples [62], multi-feature hybrids across datasets [63], recurrent networks anticipating uncertain duty [64], and signal-decomposition hybrids on noisy capacity series [65]. Transfer learning attacks generalization directly: domain-adversarial training across operating conditions [11], deep transfer with partial charging curves across cell types [12], transfer under fast-charging protocols [66], cross-chemistry EIS estimation [67], and a general framework quantifying the laboratory-to-fleet domain gap [68]. The critical assessment: accuracy claims cluster around 1–2% SOH error, but they are established overwhelmingly on public benchmark cells — 10 of the 32 studies in this complex use within-cell or k-fold validation (17 of 88 corpus-wide) — and the transfer-learning papers, the honest exception, reveal how much shift remains unaddressed. SOH estimation under controlled conditions is approaching incremental returns; calibrated lifetime and safety margins for arbitrary packs in the field remain open [2] [19].

VI. Machine Learning for Charging Control and Optimization

Charging control is the function where machine learning has changed experimental practice most directly. The defining result is Attia et al.: a Bayesian optimization loop identified high-cycle-life protocols among 224 candidates in 16 days of closed-loop cycling, against more than 500 days for an exhaustive search without early prediction [4]. Ensemble surrogates make multi-objective strategy optimization tractable [70], and Bayesian fast-charging optimization reduces charge time under degradation constraints [78]. A second line embeds degradation mechanisms into the decision itself: machine-learning fast charging that perceives and regulates internal lithium-plating onset enables plating-free rapid protocols [69], adaptive multistage schemes adjust stage boundaries against measured degradation [79], module-level optimization manages temperature gradients [80], and pack-level extensions coordinate balance-aware charging [81] [82].

The deployment gap is most concrete here: of the eight charging papers, four demonstrate on laboratory hardware — including the 224-cell closed-loop campaign above — and the remaining four rest on simulation-level evidence (three simulation-only, one with hardware spot-checks). No subset charging paper demonstrates a learned strategy validated online on a vehicle. The unsolved problem is certifiable learning-based charging — which is why mechanism-aware schemes such as plating-perception control [69] are the cluster's most deployment-relevant results.

VII. Machine Learning for Thermal Management, Faults and Safety

Thermal management (7 papers) divides by purpose. For estimation, ANN models predict cell temperature across operating conditions [71] and internal-temperature sensing extends to ultrahigh-rate discharge where surface measurement fails [83]. For design, ML surrogates amortize expensive CFD evaluation across optimization loops [84] [72], with experimentally validated parallel-liquid-cooling designs closing the loop [85]. For control, ANN controllers demonstrate parasitic-load reductions in simulation [86] — but no corpus paper closes this loop on hardware, leaving thermal control where charging control was five years ago. Physics-informed heat-generation estimation embeds thermal priors into the learning problem [9], and the cluster is consolidated by a state-of-the-art review [25].

Fault and safety applications (7 papers) divide into diagnosis of realized faults and prediction of impending hazards: grid-search-optimized SVMs classify sensor and connection faults [73], curvilinear-distance methods detect multiple pack fault types in operating fleets [87], and multi-scenario deep anomaly detection covers real EV operating envelopes [88]; force-signal deep learning catches lithium-plating-type defects at manufacturing [89]. The hazard-prediction side is increasingly model-based: multiphysics models coupled with

TABLE IV
REPRESENTATIVE CORPUS STUDIES BY MANAGEMENT FUNCTION AND METHOD FAMILY. ENTRIES GLOSS EACH STUDY AS: STUDY SUMMARY (DATA BASIS; HEADLINE CONTRIBUTION). REFERENCES IN BRACKETS REFER TO THE NUMBERED LIST.

Function	Method family	Representative studies (data; contribution)
State estimation (SOC)	Filter-embedded learning; meta-heuristic DL	LSTM+UKF embedded estimator (laboratory cells; learned observation model in a UKF) [5]; PSO-LSTM (laboratory cells; meta-heuristically tuned recurrent SOC) [32]; GA-LSTM with adaptive Kalman filtering (laboratory cells; low error under noise) [41]
Health estimation (SOH)	Classical ML; deep learning; transfer learning	ML pipeline over 179 cells (own cell fleet; relaxed charge-curve features) [28]; deep learning without extra degradation experiments (public MATR dataset; cross-chemistry generalization) [53]; real-world vehicle SOH (fleet data; field-noise robustness) [6]
Prognostics (RUL)	Classical ML; physics-informed; RL	Early-life cycle prediction (public MATR dataset; about 9% median test error) [3]; physics-informed early-cycle RUL (multiple datasets; small-data prognosis) [62]; deep-RL prognostic model fusion (public NASA dataset; adaptive model selection) [13]
Charging control	Bayesian optimization; mechanism-aware control	Closed-loop Bayesian protocol optimization (hardware prototype; 224 candidates in 16 days vs ¿500 days of exhaustive search) [4]; plating-perceiving fast charging (hardware prototype; plating-free protocols) [69]; ensemble-surrogate charging optimization (laboratory cells; multi-objective strategies) [70]
Thermal management	ANN surrogate; physics-informed; optimization	ANN temperature prediction (laboratory cells; low-error thermal model) [71]; PINN heat-generation estimation (simulation; embedded thermal physics) [9]; ML-assisted BTMS design (simulation; multi-objective cooling design) [72]
Fault & safety	Shallow ML; multiphysics + ML; transfer learning	SVM fault diagnosis (hardware prototype; sensor and connection faults) [73]; multiphysics-ML runaway prediction (simulation and hardware; module-level hazard) [74]; cross-fleet runaway-diagnosis transfer (fleet data; small-sample diagnosis) [75]
Cross-cutting	Physics-informed; XAI; digital twin; LLM	EIS diagnostics PINN review (multiple studies; paradigm synthesis) [21]; uncertainty-aware early-life prediction (public MATR dataset; calibrated forecasts) [76]; LLM-based SOH estimation with uncertainty-aware field-vehicle control (field data; corpus's earliest LLM estimator; earlier LLM-based estimation exists outside the screened venues [77]) [15]

ML predict module-level thermal-runaway progression [74], a multiphysics-informed DeepONet augments scarce experimental events with virtual data [90], and multi-source domain transfer moves thermal-runaway diagnosis across fleets with small target samples [75]. The reviews in this cluster are candid that laboratory fault diagnosis has historically failed to transfer to real vehicles [24] [23]; the fleet-data anomaly-detection studies are therefore the cluster's most valuable trend, even though label scarcity for rare faults keeps evaluation largely qualitative.

VIII. CROSS-CUTTING DIRECTIONS

A. Physics-informed learning, interpretability, and uncertainty

Physics-informed and hybrid model–data learning is the corpus's largest methodological theme (22 primary-application papers across all functions; Table I) — in state estimation [10] [91], heat-generation estimation [9], prognostics [59], and impedance-based diagnostics [21], with physics-constrained autoencoders improving health prediction over purely learned counterparts [92] and a physical-plus-data-driven fusion literature for internal-state estimation [93]. The caution is that most physics-informed studies embed simplified physics — equivalent-circuit or reduced electrochemical models — so the hybrid inherits model bias even as it gains structure, and its validation is no stronger than that of purely data-driven peers (Section III-C). A smaller but consequential cluster makes learning systems inspectable: explainability-driven feature selection improves estimator robustness [94], interpretable ML links manufacturing coating parameters to cell capacity [95], interpretable hybrid ML untangles degradation chemistries in lithium–sulfur cells [96], and uncertainty-aware models deliver calibrated early-life predictions [76]. For safety cases only calibrated uncertainty has a route into certification; attribution supports debugging.

B. Digital twins, cloud BMS, and new sensing

System-level architecture is emerging as a research object: digital twins track individual battery aging through incremental learning [97] and couple LSTM surrogates with real-time thermal observation [98]; cloud platforms perform in-situ life prediction at fleet scale [99]; second-life workflows use impedance-based accelerated SOH estimation to triage retired packs [100], and learned battery models with numerical search estimate real-time state-of-power [101]. Non-electrical sensing broadens the measurement base — ultrasonic signatures provide a SOC channel orthogonal to voltage sensing [102] — and ML mappings from equivalent-circuit parameters to electrochemical properties give classical BMS models a physically interpretable learned layer [103]. These settings differ in latency budgets, error tolerance, and update cadence, and the field is beginning to design for those differences. Production-side mapping [104], coating interpretability [95], and defect detection [89] complete the picture: cell quality is itself a BMS input.

C. Foundation models and LLM-based management

The corpus's youngest theme is generative, and it arrives with a prior-art caveat: LLM-based battery state estimation already exists outside the screened venues — preprint and conference work from 2024 and a 2025 journal letter [77]

— so no study in this corpus can claim to be the first of its kind. Within the subset the published studies are of two distinct kinds: a domain-specific LLM for risk analysis of battery *energy-storage-system* incidents — station-side rather than onboard technology — [14], and HA-LLM, which crosses the estimation boundary by using field vehicle data to estimate SOH with uncertainty-aware control, coordinating vehicle and cloud [15]. The honest framing of [15] matters for everything that follows: LLM-based estimation is no longer hypothetical — what remains normative is the claim that before any LLM-based estimator feeds actuation, it must publish calibrated intervals, structured outputs, and a latency budget, and it must not hold direct cell-level protection authority; that role belongs to physics-based envelopes and calibrated numeric estimators. Table V makes this division of labor explicit; its right-hand column is guidance proposed here, not claims about the published studies.

IX. Research Gaps, Reporting Standards, and Future Directions

Cross-referencing the bibliometric analysis of Section III with the function-level reviews yields six structural gaps. (1) Validation practice: 17 of the 88 primary papers use within-cell or unsupervised k-fold validation, and only 8 of all 103 demonstrate online or fleet use; standardized cross-cell, cross-chemistry, and lab-to-fleet transfer protocols should become reporting norms, as already practiced by the transfer-learning and general-framework studies [68] [67]. (2) Two-pole data fragmentation: the three public benchmarks sustain 23 primary papers and 36 rest on private laboratory cells that are rarely released; federated learning over fleets, anonymized fleet-data releases, and multi-institution cycling campaigns — in the spirit of the 179-cell pipeline [28] — would diversify the evidence base, and papers should at minimum disclose split definitions, seeds, and compute budgets. (3) Edge deployment: only isolated works treat onboard compute, memory, and latency as first-class constraints [58] [105]; quantization-aware training, resource-constrained architecture search, and cloud–edge hybrid designs are the natural research program. (4) Trust: uncertainty must be calibrated rather than post-hoc [76], interpretability must bound behavior rather than merely attribute it, and hybrid physics–data architectures are the most credible certification path because part of their behavior is analytically specified. (5) Foundation-model integration: an LLM-based estimator [15] and the risk-analysis LLM [14] suggest a division of labor — foundation models for reasoning and orchestration, conventional ML for calibrated estimation, physics models for safety envelopes — that should be formalized with grounding and latency guarantees. (6) Under-served functions: balancing is nearly untouched beyond pack-level insight [30], charging needs certifiable adaptive protocols, and thermal management remains simulation-dominated [25]; these are control problems where errors act on the battery, so closing them requires the safety-case machinery of gap 4 as much as new algorithms. Two follow-on subsections turn these gaps outward — first to what the automotive safety standards actually demand of these functions, then to how the emerging 2027–2030 research directions map onto them.

A. From corpus gaps to safety-standard requirements

The standards invoked in Section I — UN GTR 20 (electric-vehicle safety; Amendment 1, endorsed for WP.29 adoption in late 2026 at this review's retrieval date, adds thermal-propagation provisions with two compliance paths: a physical test, or a risk-management approach inspired by ISO 26262 and ISO 6469-1), GB 38031 (GB 38031-2020, in force through mid-2026, and the strengthened GB 38031-2025, effective July 2026, which replaces the fire/explosion-warning criterion with no-fire/no-explosion thermal-diffusion requirements and adds bottom-impact and fast-charge-cycle tests), ISO 26262:2018 (functional safety of road-vehicle E/E systems), ISO/PAS 8800:2024 (safety and artificial intelligence — the road-vehicle extension of ISO 26262 and ISO 21448 to learned components), and ISO 21448 (SOTIF) — do not prescribe estimation algorithms, but they define what a deployed estimator or controller must be able to evidence, and Table VI translates the corpus gaps into those terms. The mapping is this review's synthesis at requirement level rather than clause level, and it is stated so that it can be checked against the standards texts themselves: the point is not legal advice but that a validation protocol reported in a paper corresponds, or does not, to evidence a safety case can use. Three consequences stand out for automotive deployment. Within-cell and k-fold validation (17 primary papers) cannot support an ISO 26262 argument, because the safety case is about unseen cells and unseen duty and the protocol never tests an unseen unit; a standards argument must therefore be *argued* from the isolation evidence, not *asserted* from headline accuracy. Thermal-propagation early-warning latency is a timing requirement, so fault-detection papers evaluated on curated laboratory windows under-report both the false-alarm rate and the detection-delay distribution the GB 38031-2025 / GTR-20 escape window needs. And charging or thermal control strategies validated only in simulation lack the hardware-in-the-loop evidence base an automotive release process expects. The reporting checklist below (Table VIII) is deliberately aligned with this mapping — reporting an isolation argument and a validation environment is the minimum on which any standards conversation can start.

B. Position relative to prior reviews

The closest published competitor is a 2026 performance-centric review of ML for EV powertrains with BMS and charging scope [27]; adjacent cross-functional surveys exist for degradation and estimation methods. This review is distinguished by what those studies do not attempt, and Table VII states the comparison: a pooled two-tier corpus with record-level screening provenance, a within-corpus measurement of validation protocol strength across functions, a standards-grounded gap analysis, and a released, re-run-able analysis pipeline. Where prior reviews ask "which method works best for function X," this review asks "which claims can actually be checked — and what would it cost the reader to trust them," which is the question a scoping review with a curated classified subset is positioned to answer.

TABLE V
SAFETY AND VERIFICATION TAXONOMY FOR FOUNDATION-MODEL COMPONENTS IN THE BMS STACK. THE RIGHT-HAND COLUMN IS NORMATIVE GUIDANCE PROPOSED BY THIS REVIEW, NOT CLAIMS ABOUT THE PUBLISHED STUDIES.

Dimension	**Classical / deep ML estimators (established practice)**	**Published LLM studies (2025–2026)**	**Required guardrail (this review)**
Role in the stack	Calibrated numeric estimation and detection	Stationary-BESS incident risk analysis [14]; LLM-based SOH estimation with uncertainty-aware control on field vehicle data [15]	Estimation should belong to calibrated ML; the LLM SOH estimator of [15] is the earliest in the screened subset — LLM-based battery estimation exists earlier outside these venues [77] — and must show calibrated intervals before any actuation-adjacent role
Grounding	Trained on domain cycling data	Domain-tuned corpora and knowledge-enhanced prompting	Retrieval over verified corpora with source citations; no ungrounded recall
Dominant failure mode	Bounded numeric error under distribution shift	Open-ended hallucinated text	Outputs restricted to structured, checkable forms; never free text into protection logic
Uncertainty	Gaussian-process / conformal intervals in a subset of studies	Not reported for incident risk analysis [14]; uncertainty-aware weighting demonstrated in the SOH estimator [15]	Calibrated intervals mandatory before any actuation-adjacent role
Latency and compute	Millisecond-class, onboard	Second-class, cloud-dependent	Offline analysis and planning only unless real-time evidence is shown
Safety authority	Bounded by physics envelopes and thresholds	Undefined in published studies	LLM components must never hold direct cell-level protection authority

TABLE VI
MAPPING BMS FUNCTIONS TO REQUIREMENTS IN UN GTR 20 (AMENDMENT 1 THERMAL-PROPAGATION PROVISIONS), GB 38031-2020/2025, ISO 26262:2018 WITH ISO/PAS 8800:2024, AND ISO 21448 (SOTIF): WHAT DEPLOYED ML COMPONENTS MUST BE ABLE TO EVIDENCE. THE MAPPING IS THIS REVIEW'S SYNTHESIS AT REQUIREMENT LEVEL (NOT CLAUSE-LEVEL LEGAL CITATION); "CORPUS STATUS" GIVES THE CORRESPONDING CLASSIFIED-SUBSET PRIMARY-PAPER COUNTS FROM SECTION III-C.

BMS function	**What the standard demands of a deployed function**	**Current corpus status**
SOC estimation	ISO 26262:2018 treats estimation error that can enable overcharge/overdischarge as an E/E-system safety hazard to be handled through the safety lifecycle; ISO/PAS 8800:2024 extends this to learned elements (assurance argument over data, model, and known insufficiencies); UN GTR 20 / GB 38031-2025 treat overcharge as an abuse-protection requirement	3 of 13 SOC primary papers use a strong isolation protocol; none validate on real-vehicle data; 12 rest on laboratory-scale data
SOH / RUL prognosis	Capacity and power fade margins underpin thermal-propagation early warning and end-of-service decisions; degradation assumptions must hold across chemistry, duty, and environment (ISO 26262 hazard analysis; ISO 21448 for performance limitations)	20 of 32 SOH/RUL primary papers use a strong protocol; 4 use real-vehicle data; accuracy claims rest mostly on three public benchmarks (23 papers corpus-wide)
Fault detection	Thermal-propagation early warning must provide the escape window: GB 38031-2025 requires the warning signal but no fire/explosion beyond the triggered cell (superseding the GB 38031-2020 5-min criterion); the GTR 20 amendment proposes a $\geq$5-min advance warning before cabin hazards, via test or risk-management compliance path; coverage must be argued and audited (ISO/PAS 8800 assurance argument)	2 of 7 fault/safety primary papers validate online on fleet data; detection-latency and false-alarm distributions are rarely reported
Charging control	Overcharge, over-temperature, and thermal-cycling limits are hard constraints under GB 38031-2020/2025; adaptive strategies must be bounded by protection envelopes and argued in the safety case	0 of 8 charging primary papers use a strong protocol; 4 demonstrate on laboratory hardware, none online
Thermal management	Temperature-gradient and propagation-containment requirements after fast-charge or abuse (GB 38031-2025 thermal diffusion, bottom-impact, fast-charge-cycle tests); the cooling function must be verified on hardware	0 of 7 thermal primary papers use a strong protocol; 2 exercise laboratory hardware; no learned control loop is validated on hardware

TABLE VII
POSITION RELATIVE TO PRIOR REVIEWS. "PRASANNA 2026" DENOTES A 2026 PERFORMANCE-CENTRIC ML REVIEW OF EV POWERTRAINS WITH BMS AND CHARGING SCOPE; FUNCTION- AND METHOD-SPECIFIC REVIEWS ARE SUMMARIZED COLLECTIVELY. THIS REVIEW DIFFERS MOST IN THE VALIDATION MEASUREMENT (WITHIN A CURATED SUBSET), STANDARDS LINKAGE, AND RELEASED RECORD-LEVEL ARTIFACTS.

Axis	**This review**	**Prior cross-functional / powertrain reviews**
Corpus	Two-tier: 799 rule-based eligible records + 103 classified subset	Fixed study lists (e.g., 95 studies), no record-level selection provenance
Validation analysis	Protocol classes defined, per-paper assignments released, method-family contrast tested with sensitivity ranges; within-corpus measurement, not a field census	Narrative remarks on data scarcity; protocol strength not quantified
Standards linkage	Function-to-requirement mapping vs. UN GTR 20, GB 38031-2020/2025, ISO 26262:2018, ISO/PAS 8800:2024	Not addressed or addressed qualitatively
Recommendations	Twelve-item reporting checklist (BMS-ML-RC) + six research directions	General future-work statements
Reproducibility	Released papers.csv, screening provenance, codebook, and statistics scripts	No released record-level dataset

C. Outlook: what the next revision cycle of this literature should produce

Concretely, and tied to the gaps above, the highest-value near-term work falls in six directions. (a) *Duty-cycle-aware validation*: public cross-cell benchmarks that pair cycling data with recorded EV drive profiles, so that cross-chemistry and cross-domain generalization can be tested against transportation duty rather than laboratory schedules; (b) *federated fleet evidence*: privacy-preserving SOH/SOC studies over live fleets (≈2% of the eligible corpus today), with the 179-cell pipeline as the template; (c) *benchmark diversification*: multi-institution cycling campaigns two orders of magnitude larger than MATR/NASA/CALCE and released with DOI-backed artifacts, addressing the two-pole fragmentation of Section III-C; (d) *pack- and vehicle-level ML*: estimation validated at the integration level where UN GTR 20 and GB 38031 requirements bite, instead of cell-level evidence transplanted upward; (e) *functionally safe learning*: ML estimators whose uncertainty handling and assurance arguments are specified against ISO 26262 safety goals and ISO/PAS 8800, with ISO 21448/SOTIF reasoning for capability limitations, making certifiable charging and thermal control reachable; and (f) *honest foundation-model integration*: retrieval-grounded LLM agents with published latency and calibration budgets, capable of the coordination role while estimation authority stays with calibrated numeric models — extending, not contradicting, the corpus's LLM estimator [15]. A corpus-level guardrail on all six: the BMS-ML-RC checklist below is designed to be attached to a paper at submission, so the next revision cycle of this literature arrives auditable by construction.

D. A reporting checklist for machine-learning BMS studies

The most concrete artifact this review can offer is procedural. Modeled on the CONSORT and PRISMA checklists, Table VIII proposes a twelve-item BMS-ML reporting checklist (BMS-ML-RC): the minimum metadata a study should disclose for its accuracy claims to be comparable, reproducible, and poolable — and, per Section IX-A, admissible in a safety discussion. Its derivation is explicit: items 4, 5, and 11 operationalize the validation-protocol classes defined in Section II-C and measured in the classified subset; items 1–3 and 6–10 respond to the reporting omissions observed across Sections III–VIII; item 12 is the artifact-availability norm this paper itself follows. What the checklist is and is not: it is a proposal grounded in the classified subset's measured weaknesses, not a community-validated instrument — no Delphi or expert-consensus process was run for it, only three items are score-measured in this corpus, and it has not been applied as an audit to all 103 studies. The review's own reporting follows the PRISMA-ScR flow structure (Fig. 2), but the protocol was not prospectively registered; where this section says a practice *should* be reported, that is a recommendation supported by the corpus evidence above, and we suggest journals consider it as a submission checklist.

TABLE VIII
PROPOSED BMS-ML REPORTING CHECKLIST (BMS-ML-RC), MODELED ON CONSORT/PRISMA REPORTING STANDARDS. ITEMS 4, 5, AND 11 DERIVE FROM THE VALIDATION-PROTOCOL CLASSES DEFINED IN SECTION II-C AND MEASURED IN THE CLASSIFIED SUBSET; ITEM 12 IS THE ARTIFACT-AVAILABILITY NORM THIS PAPER ITSELF FOLLOWS. THE CHECKLIST AS A WHOLE HAS NOT BEEN SCORE-AUDITED ACROSS ALL 103 CLASSIFIED PAPERS — IT IS A PROPOSAL, NOT A SURVEY RESULT.

#	Reporting item — what must be reported
1	Dataset identity — exact name and version of every dataset (e.g., Severson-MATR, NASA, CALCE); download source
2	Cell population — number of cells, chemistry, manufacturer, format, and batch for every training and test set
3	Sensor and sampling context — measured channels, sampling rate, sensor class, temperature conditions
4	Split protocol — split type (within-cell / cross-cell / cross-chemistry / cross-domain) and isolation argument
5	Out-of-distribution evaluation — at least one protocol where test conditions differ from training
6	Baselines — at least one classical baseline alongside the proposed architecture, on identical data
7	Randomness control — seeds or repeated-run statistics for every stochastic result
8	Uncertainty — calibrated intervals or error bars for headline claims, not point estimates alone
9	Compute budget — training compute and inference latency/memory on the target platform
10	Failure analysis — at least one documented failure mode and its distribution
11	Validation environment — simulation, laboratory cycling, hardware prototype, or online/fleet
12	Artifact availability — data and code deposit (DOI-backed) or an explicit stated exemption

X. CONCLUSION

This review has examined machine learning in battery management on a two-tier evidence base: a rule-based eligible corpus of 799 journal articles from which field-level shares are measured within a documented capped retrieval, and an in-depth classified subset of 103 studies through which method, data, and validation practice are examined with a 30-paper blind reliability check and label-noise sensitivity ranges. Three conclusions stand out. First, within the classified subset, method families have shifted from classical feature-based learners to deep sequence models and now physics-informed hybrids, with LLM-based studies appearing from 2025–2026 — though LLM-based battery estimation began in other venues. Second, the center of gravity is estimation at cell level: measured as shares of the eligible corpus, the SOC–SOH–prognostics union covers about half of function-tagged records, whereas vehicle- and pack-level practice remains thin — precisely the level at which UN GTR 20, GB 38031, and ISO 26262 pose their requirements. Third, published accuracy progress is not matched by validation progress: only 8 of 103 classified studies demonstrate online or fleet operation, 17 of 88 primary papers use isolation-weakened splits, and 23 rest on three public benchmarks while 36 rest on private laboratory cells; the physics-informed validation contrast is directionally consistent under both reported definitions but not statistically established (protocol-only 36% versus 52%, $p = 0.32$; including-hardware 59% versus 74%, $p = 0.19$), so it should be treated as an untested hypothesis rather than a settled result. The near-term agenda follows directly: standardized cross-domain validation, diversified fleet data, pack- and vehicle-level ML aligned with the standards mapping of Section IX, calibrated uncertainty, and a strict division of labor between foundation models, calibrated estimators, and

physics-based safety envelopes. If the last five years were about discovering what machine learning can estimate about a battery, the next five will be decided by what it can be trusted to do — onboard, at fleet scale, across chemistries, and over full cell lifetimes.

Data and Code Availability

The complete evidentiary record of this review is released as a Zenodo deposit (DOI: https://doi.org/10.5281/zenodo.23436202): the classified corpus of 103 studies (papers.csv) with per-paper citation counts, the 799-record eligible corpus and its record-level screening provenance (screening_log.csv, all 1,550 screening decisions plus five forced-seed entries), the released title-gate and tagging term lists (tagging_terms.json), the per-paper validation-protocol audit (protocol_audit.csv), the classification codebook, the blind 30-paper inter-coder reliability study (locked coding sheets, agreement report, itemized disagreement table), and the label-noise sensitivity analysis over headline counts. The complete search, screening, verification, and figure-generation scripts ship inside the same deposit and regenerate every figure and count in Section III; a development-form mirror is kept on GitHub.

Acknowledgment

The authors would like to acknowledge the use of writing-support and language-enhancement tools, including OpenAI ChatGPT, Grammarly, and QuillBot, during the preparation of this manuscript. These tools were used to assist with grammar correction, readability improvement, language refinement, and sentence restructuring. All technical content, analysis, interpretations, and research contributions are from the authors themselves. This work received no specific funding.

References

[1] M. F. Ng, J. Zhao, Q. Yan, G. J. Conduit, Z. W. Seh, "Predicting the state of charge and health of batteries using data-driven machine learning," *Nature Machine Intelligence*, vol. 2, no. 3, p. 161-170, 2020. doi: 10.1038/s42256-020-0156-7.

[2] X. Hu, L. Xu, X. Lin, M. Pecht, "Battery Lifetime Prognostics," *Joule*, vol. 4, no. 2, p. 310-346, 2020. doi: 10.1016/j.joule.2019.11.018.

[3] K. A. Severson, P. M. Attia, N. Jin, N. Perkins, B. Jiang, Z. Yang et al., "Data-driven prediction of battery cycle life before capacity degradation," *Nature Energy*, vol. 4, no. 5, p. 383-391, 2019. doi: 10.1038/s41560-019-0356-8.

[4] P. M. Attia, A. Grover, N. Jin, K. A. Severson, T. M. Markov, Y. H. Liao et al., "Closed-loop optimization of fast-charging protocols for batteries with machine learning," *Nature*, vol. 578, no. 7795, p. 397-402, 2020. doi: 10.1038/s41586-020-1994-5.

[5] F. Yang, S. Zhang, W. Li, Q. Miao, "State-of-charge estimation of lithium-ion batteries using LSTM and UKF," *Energy*, vol. 201, p. 117664, 2020. doi: 10.1016/j.energy.2020.117664.

[6] J. Tian, X. Liu, S. Li, Z. Wei, X. Zhang, G. Xiao et al., "Lithium-ion battery health estimation with real-world data for electric vehicles," *Energy*, vol. 270, p. 126855, 2023. doi: 10.1016/j.energy.2023.126855.

[7] Y. Wei, D. Wu, "Prediction of state of health and remaining useful life of lithium-ion battery using graph convolutional network with dual attention mechanisms," *Reliability Engineering & System Safety*, vol. 230, p. 108947, 2023. doi: 10.1016/j.ress.2022.108947.

[8] Z. Li, X. Zhang, W. Gao, "State of health estimation of lithium-ion battery during fast charging process based on BiLSTM-Transformer," *Energy*, vol. 311, p. 133418, 2024. doi: 10.1016/j.energy.2024.133418.

[9] H. Pang, L. Wu, J. Liu, X. Liu, K. Liu, "Physics-informed neural network approach for heat generation rate estimation of lithium-ion battery under various driving conditions," *Journal of Energy Chemistry*, vol. 78, p. 1-12, 2023. doi: 10.1016/j.jechem.2022.11.036.

[10] Y. Wang, X. Han, D. Guo, L. Lu, Y. Chen, M. Ouyang, "Physics-Informed Recurrent Neural Network With Fractional-Order Gradients for State-of-Charge Estimation of Lithium-Ion Battery," *IEEE Journal of Radio Frequency Identification*, vol. 6, p. 968-971, 2022. doi: 10.1109/jrfid.2022.3211841.

[11] Z. Ye, J. Yu, "State-of-Health Estimation for Lithium-Ion Batteries Using Domain Adversarial Transfer Learning," *IEEE Transactions on Power Electronics*, vol. 37, no. 3, p. 3528-3543, 2022. doi: 10.1109/tpel.2021.3117788.

[12] J. Yao, T. Han, "Data-driven lithium-ion batteries capacity estimation based on deep transfer learning using partial segment of charging/discharging data," *Energy*, vol. 271, p. 127033, 2023. doi: 10.1016/j.energy.2023.127033.

[13] M. Du, Y. Zhang, Q. Miao, "Remaining useful life prediction of lithium battery based on deep reinforcement learning fusion network," *Reliability Engineering & System Safety*, vol. 264, p. 111392, 2025. doi: 10.1016/j.ress.2025.111392.

[14] J. Liang, H. Meng, Y. Mu, "Domain-specific large language model-driven risk analysis of battery energy storage systems," *Reliability Engineering & System Safety*, vol. 274, p. 112416, 2026. doi: 10.1016/j.ress.2026.112416.

[15] S. Cai, J. Chen, J. Hu, R. Huang, J. Liao, M. Wei et al., "HA-LLM: A domain knowledge-enhanced large language model for vehicle-cloud collaborative battery health awareness," *Energy Storage Materials*, vol. 86, p. 104983, 2026. doi: 10.1016/j.ensm.2026.104983.

[16] Y. Liu, Y. He, H. Bian, W. Guo, X. Zhang, "A review of lithium-ion battery state of charge estimation based on deep learning: Directions for improvement and future trends," *Journal of Energy Storage*, vol. 52, p. 104664, 2022. doi: 10.1016/j.est.2022.104664.

[17] Y. Wu, D. Bai, K. Zhang, Y. Li, F. Yang, "Advancements in the estimation of the state of charge of lithium-ion battery: a comprehensive review of traditional and deep learning approaches," *Journal of Materials Informatics*, vol. 5, no. 2, 2025. doi: 10.20517/jmi.2024.84.

[18] X. Sui, S. He, S. B. Vilsen, J. Meng, R. Teodorescu, D. I. Stroe, "A review of non-probabilistic machine learning-based state of health estimation techniques for Lithium-ion battery," *Applied Energy*, vol. 300, p. 117346, 2021. doi: 10.1016/j.apenergy.2021.117346.

[19] X. Li, D. Yu, V. Søren Byg, S. Daniel Ioan, "The development of machine learning-based remaining useful life prediction for lithium-ion batteries," *Journal of Energy Chemistry*, vol. 82, p. 103-121, 2023. doi: 10.1016/j.jechem.2023.03.026.

[20] J. Yao, J. Kowal, "Towards a smarter battery management system: A critical review on deep learning-based state of charge estimation of lithium-ion batteries," *Energy and AI*, vol. 21, p. 100585, 2025. doi: 10.1016/j.egyai.2025.100585.

[21] Y. H. Lin, S. J. Ruan, Y. X. Chen, Y. F. Li, "Physics-informed deep learning for lithium-ion battery diagnostics using electrochemical impedance spectroscopy," *Renewable and Sustainable Energy Reviews*, vol. 188, p. 113807, 2023. doi: 10.1016/j.rser.2023.113807.

[22] T. Wang, K. Liu, H. Peng, X. Liu, "Interpretable Machine Learning for Battery Prognosis: Retrospect and Prospect," *Advanced Energy Materials*, vol. 15, no. 48, Art. no. e03067, 2025. doi: 10.1002/aenm.202503067.

[23] A. Samanta, S. Chowdhuri, S. S. Williamson, "Machine Learning-Based Data-Driven Fault Detection/Diagnosis of Lithium-Ion Battery: A Critical Review," *Electronics*, vol. 10, no. 11, p. 1309, 2021. doi: 10.3390/electronics10111309.

[24] Q. Yu, C. Wang, J. Li, R. Xiong, M. Pecht, "Challenges and outlook for lithium-ion battery fault diagnosis methods from the laboratory to real world applications," *eTransportation*, vol. 17, p. 100254, 2023. doi: 10.1016/j.etran.2023.100254.

[25] A. Li, J. Weng, A. C. Y. Yuen, W. Wang, H. Liu, E. W. M. Lee et al., "Machine learning assisted advanced battery thermal management system: A state-of-the-art review," *Journal of Energy Storage*, vol. 60, p. 106688, 2023. doi: 10.1016/j.est.2023.106688.

[26] T. Rahman, T. Alharbi, "Exploring Lithium-Ion Battery Degradation: A Concise Review of Critical Factors, Impacts, Data-Driven Degradation Estimation Techniques, and Sustainable Directions for Energy Storage Systems," *Batteries*, vol. 10, no. 7, p. 220, 2024. doi: 10.3390/batteries10070220.

[27] R. Prasanna, R. S. Kumar, N. S. Bhuvaneswari, C. Karthik, "Performance centric review of machine learning techniques for electric vehicle powertrain with battery management and charging systems," *Discover Artificial Intelligence*, vol. 6, no. 1, Art. no. 141, 2026. doi: 10.1007/s44163-025-00721-y.

[28] D. Roman, S. Saxena, V. Robu, M. Pecht, D. Flynn, "Machine learning pipeline for battery state-of-health estimation," *Nature Machine Intelligence*, vol. 3, no. 5, p. 447-456, 2021. doi: 10.1038/s42256-021-00312-3.

[29] N. M. Salem, A. Mohamed, "Physics-informed hybrid reinforcement learning for estimating lithium-ion battery state of health," *Scientific Reports*, vol. 15, no. 1, Art. no. 43819, 2025. doi: 10.1038/s41598-025-30602-4.

[30] J. Hardy, J. Steggall, P. Hardy, "Rethinking lithium-ion battery management: Eliminating routine cell balancing enhances hazardous fault detection," *Journal of Energy Storage*, vol. 63, p. 106931, 2023. doi: 10.1016/j.est.2023.106931.

[31] Z. Deng, X. Hu, X. Lin, Y. Che, L. Xu, W. Guo, "Data-driven state of charge estimation for lithium-ion battery packs based on Gaussian process regression," *Energy*, vol. 205, p. 118000, 2020. doi: 10.1016/j.energy.2020.118000.

[32] X. Ren, S. Liu, X. Yu, X. Dong, "A method for state-of-charge estimation of lithium-ion batteries based on PSO-LSTM," *Energy*, vol. 234, p. 121236, 2021. doi: 10.1016/j.energy.2021.121236.

[33] M. K. Khan, M. A. Houran, K. Kauhaniemi, M. H. Zafar, M. Mansoor, S. Rashid, "Efficient state of charge estimation of lithium-ion batteries in electric vehicles using evolutionary intelligence-assisted GLA–CNN–Bi-LSTM deep learning model," *Heliyon*, vol. 10, no. 15, p. e35183, 2024. doi: 10.1016/j.heliyon.2024.e35183.

[34] M. Ragone, V. Yurkiv, A. Ramasubramanian, B. Kashir, F. Mashayek, "Data driven estimation of electric vehicle battery state-of-charge informed by automotive simulations and multi-physics modeling," *Journal of Power Sources*, vol. 483, p. 229108, 2021. doi: 10.1016/j.jpowsour.2020.229108.

[35] V. Chandran, C. K. Patil, A. Karthick, D. Ganeshaperumal, R. Rahim, A. Ghosh, "State of Charge Estimation of Lithium-Ion Battery for Electric Vehicles Using Machine Learning Algorithms," *World Electric Vehicle Journal*, vol. 12, no. 1, p. 38, 2021. doi: 10.3390/wevj12010038.

[36] M. Naguib, P. Kollmeyer, A. Emadi, "Lithium-Ion Battery Pack Robust State of Charge Estimation, Cell Inconsistency, and Balancing: Review," *IEEE Access*, vol. 9, p. 50570-50582, 2021. doi: 10.1109/access.2021.3068776.

[37] J. Tian, R. Xiong, W. Shen, J. Lu, F. Sun, "Flexible battery state of health and state of charge estimation using partial charging data and deep learning," *Energy Storage Materials*, vol. 51, p. 372-381, 2022. doi: 10.1016/j.ensm.2022.06.053.

[38] M. Zou, J. Wang, D. Yan, Y. Li, X. Tang, "Novel state of charge estimation method of containerized Lithium–Ion battery energy storage system based on deep learning," *Journal of Power Sources*, vol. 624, p. 235609, 2024. doi: 10.1016/j.jpowsour.2024.235609.

[39] S. Arandhakar, J. Nakka, "State of charge estimation of lithium ion battery for electric vehicle using cutting edge machine learning algorithms: A review," *Journal of Energy Storage*, vol. 103, p. 114281, 2024. doi: 10.1016/j.est.2024.114281.

[40] U. Khan, S. Kirmani, Y. Rafat, M. U. Rehman, M. S. Alam, "Improved deep learning based state of charge estimation of lithium ion battery for electrified transportation," *Journal of Energy Storage*, vol. 91, p. 111877, 2024. doi: 10.1016/j.est.2024.111877.

[41] J. Wang, Z. Zuo, Y. Wei, Y. Jia, B. Chen, Y. Li et al., "State of charge estimation of lithium-ion battery based on GA-LSTM and improved IAKF," *Applied Energy*, vol. 368, p. 123508, 2024. doi: 10.1016/j.apenergy.2024.123508.

[42] E. Buchicchio, A. De Angelis, F. Santoni, P. Carbone, F. Bianconi, F. Smeraldi, "Battery SOC estimation from EIS data based on machine learning and equivalent circuit model," *Energy*, vol. 283, p. 128461, 2023. doi: 10.1016/j.energy.2023.128461.

[43] X. Li, C. Yuan, X. Li, Z. Wang, "State of health estimation for Li-Ion battery using incremental capacity analysis and Gaussian process regression," *Energy*, vol. 190, p. 116467, 2020. doi: 10.1016/j.energy.2019.116467.

[44] Z. Wang, C. Yuan, X. Li, "Lithium Battery State-of-Health Estimation via Differential Thermal Voltammetry With Gaussian Process Regression," *IEEE Transactions on Transportation Electrification*, vol. 7, no. 1, p. 16-25, 2021. doi: 10.1109/tte.2020.3028784.

[45] Y. Zhang, Q. Tang, Y. Zhang, J. Wang, U. Stimming, A. A. Lee, "Identifying degradation patterns of lithium ion batteries from impedance spectroscopy using machine learning," *Nature Communications*, vol. 11, no. 1, Art. no. 1706, 2020. doi: 10.1038/s41467-020-15235-7.

[46] P. Tagade, K. S. Hariharan, S. Ramachandran, A. Khandelwal, A. Naha, S. M. Kolake et al., "Deep Gaussian process regression for lithium-ion battery health prognosis and degradation mode diagnosis," *Journal of Power Sources*, vol. 445, p. 227281, 2020. doi: 10.1016/j.jpowsour.2019.227281.

[47] Z. Deng, X. Hu, P. Li, X. Lin, X. Bian, "Data-Driven Battery State of Health Estimation Based on Random Partial Charging Data," *IEEE Transactions on Power Electronics*, vol. 37, no. 5, p. 5021-5031, 2022. doi: 10.1109/tpel.2021.3134701.

[48] K. Zheng, J. Meng, Z. Yang, F. Zhou, K. Yang, Z. Song, "Refined lithium-ion battery state of health estimation with charging segment adjustment," *Applied Energy*, vol. 375, p. 124077, 2024. doi: 10.1016/j.apenergy.2024.124077.

[49] K. Q. Zhou, Y. Qin, C. Yuen, "Graph neural network-based lithium-ion battery state of health estimation using partial discharging curve," *Journal of Energy Storage*, vol. 100, p. 113502, 2024. doi: 10.1016/j.est.2024.113502.

[50] J. Zhao, X. Qu, Y. Li, J. Nan, A. F. Burke, "Real-time prediction of battery remaining useful life using hybrid-fusion deep neural networks," *Energy*, vol. 328, p. 136618, 2025. doi: 10.1016/j.energy.2025.136618.

[51] B. Zhao, W. Zhang, Y. Zhang, C. Zhang, C. Zhang, J. Zhang, "Lithium-ion battery remaining useful life prediction based on interpretable deep learning and network parameter optimization," *Applied Energy*, vol. 379, p. 124713, 2025. doi: 10.1016/j.apenergy.2024.124713.

[52] X. Qiang, W. Liu, Z. Lyu, H. Ruan, X. Li, "A data-fusion-model method for state of health estimation of Li-ion battery packs based on partial charging curve," *Green Energy and Intelligent Transportation*, vol. 3, no. 5, p. 100169, 2024. doi: 10.1016/j.geits.2024.100169.

[53] J. Lu, R. Xiong, J. Tian, C. Wang, F. Sun, "Deep learning to estimate lithium-ion battery state of health without additional degradation experiments," *Nature Communications*, vol. 14, no. 1, Art. no. 2760, 2023. doi: 10.1038/s41467-023-38458-w.

[54] Y. Yang, "A machine-learning prediction method of lithium-ion battery life based on charge process for different applications," *Applied Energy*, vol. 292, p. 116897, 2021. doi: 10.1016/j.apenergy.2021.116897.

[55] Z. Fei, F. Yang, K. L. Tsui, L. Li, Z. Zhang, "Early prediction of battery lifetime via a machine learning based framework," *Energy*, vol. 225, p. 120205, 2021. doi: 10.1016/j.energy.2021.120205.

[56] Z. Tong, J. Miao, S. Tong, Y. Lu, "Early prediction of remaining useful life for Lithium-ion batteries based on a hybrid machine learning method," *Journal of Cleaner Production*, vol. 317, p. 128265, 2021. doi: 10.1016/j.jclepro.2021.128265.

[57] Q. Si, S. Matsuda, Y. Yamaji, T. Momma, Y. Tateyama, "Data-Driven Cycle Life Prediction of Lithium Metal-Based Rechargeable Battery Based on Discharge/Charge Capacity and Relaxation Features," *Advanced Science*, vol. 11, no. 33, Art. no. 2402608, 2024. doi: 10.1002/advs.202402608.

[58] S. Zhao, C. Zhang, Y. Wang, "Lithium-ion battery capacity and remaining useful life prediction using board learning system and long short-term memory neural network," *Journal of Energy Storage*, vol. 52, p. 104901, 2022. doi: 10.1016/j.est.2022.104901.

[59] J. Shi, A. Rivera, D. Wu, "Battery health management using physics-informed machine learning: Online degradation modeling and remaining useful life prediction," *Mechanical Systems and Signal Processing*, vol. 179, p. 109347, 2022. doi: 10.1016/j.ymssp.2022.109347.

[60] S. Wang, R. Zhou, Y. Ren, H. Liu, Y. Lin, C. Lian, "A generalizable physics-informed neural network for lithium-ion battery SOH estimation utilizing partial charging segments," *Journal of Energy Chemistry*, vol. 112, p. 977-986, 2026. doi: 10.1016/j.jechem.2025.08.093.

[61] K. Gao, Q. Li, L. Hu, J. Huang, H. Li, Y. Wu, "Physical informed neural network for SOH estimation of lithium-ion battery with electrochemical mechanism," *Energy*, vol. 342, p. 139753, 2026. doi: 10.1016/j.energy.2025.139753.

[62] W. Deng, H. Le, K. T. Nguyen, C. Gogu, K. Medjaher, J. Morio et al., "A generic physics-informed machine learning framework for battery remaining useful life prediction using small early-stage lifecycle data," *Applied Energy*, vol. 384, p. 125314, 2025. doi: 10.1016/j.apenergy.2025.125314.

[63] Y. Qi, J. Wang, H. Lu, H. Chu, D. Chen, J. Jin, "Remaining useful life prediction for lithium-ion battery based on hybrid machine learning with Wiener process," *Journal of Energy Storage*, vol. 136, p. 118415, 2025. doi: 10.1016/j.est.2025.118415.

[64] J. Lu, R. Xiong, J. Tian, C. Wang, C. W. Hsu, N. T. Tsou et al., "Battery degradation prediction against uncertain future conditions with recurrent neural network enabled deep learning," *Energy Storage Materials*, vol. 50, p. 139-151, 2022. doi: 10.1016/j.ensm.2022.05.007.

[65] Y. Fan, Z. Lin, F. Wang, J. Zhang, "A hybrid approach for lithium-ion battery remaining useful life prediction using signal decomposition and machine learning," *Scientific Reports*, vol. 15, no. 1, Art. no. 8161, 2025. doi: 10.1038/s41598-025-92262-8.

[66] J. Zhao, D. Li, Y. Li, D. Shi, J. Nan, A. F. Burke, "Battery state of health estimation under fast charging via deep transfer learning," *iScience*, vol. 28, no. 5, p. 112235, 2025. doi: 10.1016/j.isci.2025.112235.

[67] M. Lin, D. Hu, J. Meng, J. Wu, "Transfer Learning-Based Lithium-Ion Battery State of Health Estimation With Electrochemical Impedance Spectroscopy," *IEEE Transactions on Transportation Electrification*, vol. 11, no. 3, p. 7910-7920, 2025. doi: 10.1109/tte.2025.3533540.

[68] Z. Cao, W. Gao, Y. Fu, N. V. Kurdkandi, C. Mi, "A general framework for lithium-ion battery state of health estimation: From laboratory tests to machine learning with transferability across domains," *Applied Energy*, vol. 381, p. 125086, 2025. doi: 10.1016/j.apenergy.2024.125086.

[69] Z. Wei, X. Yang, Y. Li, H. He, W. Li, D. U. Sauer, "Machine learning-based fast charging of lithium-ion battery by perceiving and regulating internal microscopic states," *Energy Storage Materials*, vol. 56, p. 62-75, 2023. doi: 10.1016/j.ensm.2022.12.034.

[70] F. Cheng, H. Liu, "Charging strategies optimization for lithium-ion battery: Heterogeneous ensemble surrogate model-assisted advanced multi-objective optimization algorithm," *Energy Conversion and Management*, vol. 342, p. 120170, 2025. doi: 10.1016/j.enconman.2025.120170.

[71] Y. Wang, X. Chen, C. Li, Y. Yu, G. Zhou, C. Wang et al., "Temperature prediction of lithium-ion battery based on artificial neural network model," *Applied Thermal Engineering*, vol. 228, p. 120482, 2023. doi: 10.1016/j.applthermaleng.2023.120482.

[72] X. Zhou, W. Guo, X. Shi, C. She, Z. Zheng, J. Zhou et al., "Machine learning assisted multi-objective design optimization for battery thermal management system," *Applied Thermal Engineering*, vol. 253, p. 123826, 2024. doi: 10.1016/j.applthermaleng.2024.123826.

[73] L. Yao, Z. Fang, Y. Xiao, J. Hou, Z. Fu, "An Intelligent Fault Diagnosis Method for Lithium Battery Systems Based on Grid Search Support Vector Machine," *Energy*, vol. 214, p. 118866, 2021. doi: 10.1016/j.energy.2020.118866.

[74] B. R. Das Goswami, Y. Abdisobbouhi, H. Du, F. Mashayek, T. A. Kingston, V. Yurkiv, "Advancing battery safety: Integrating multiphysics and machine learning for thermal runaway prediction in lithium-ion battery module," *Journal of Power Sources*, vol. 614, p. 235015, 2024. doi: 10.1016/j.jpowsour.2024.235015.

[75] C. Dong, D. Sun, "Multi-source domain transfer learning with small sample learning for thermal runaway diagnosis of lithium-ion battery," *Applied Energy*, vol. 365, p. 123248, 2024. doi: 10.1016/j.apenergy.2024.123248.

[76] L. H. Rieger, E. Flores, K. F. Nielsen, P. Norby, E. Ayerbe, O. Winther et al., "Uncertainty-aware and explainable machine learning for early prediction of battery degradation trajectory," *Digital Discovery*, vol. 2, no. 1, p. 112-122, 2023. doi: 10.1039/d2dd00067a.

[77] W. Q. T. Poh, Y. Xu, R. T. P. Tan, "Data-Driven Estimation of Li-Ion Battery Health: Integrating GPT-4 With Distilled Lifelong Learning," *IEEE Transactions on Energy Conversion*, vol. 40, no. 2, p. 1682-1685, 2025. doi: 10.1109/tec.2025.3548400.

[78] G. Dong, Y. Feng, Y. Lou, M. Zhang, J. Wei, "Data-Driven Fast Charging Optimization for Lithium-Ion Battery Using Bayesian Optimization With Fast Convergence," *IEEE Transactions on Transportation Electrification*, vol. 10, no. 2, p. 4173-4183, 2024. doi: 10.1109/tte.2023.3311813.

[79] Y. Li, K. Li, Y. Xie, J. Liu, C. Fu, B. Liu, "Optimized charging of lithium-ion battery for electric vehicles: Adaptive multistage constant current–constant voltage charging strategy," *Renewable Energy*, vol. 146, p. 2688-2699, 2020. doi: 10.1016/j.renene.2019.08.077.

[80] Z. Fan, Y. Fu, H. Liang, R. Gao, S. Liu, "A module-level charging optimization method of lithium-ion battery considering temperature gradient effect of liquid cooling and charging time," *Energy*, vol. 265, p. 126331, 2023. doi: 10.1016/j.energy.2022.126331.

[81] A. Heydarian Ardakani, F. Abdollahi, "A fast balance optimization approach for charging enhancement of lithium-ion battery packs through deep reinforcement learning," *Journal of Energy Storage*, vol. 89, p. 111755, 2024. doi: 10.1016/j.est.2024.111755.

[82] S. Duan, K. Xia, J. Li, Z. Zhao, H. Liu, "Optimization charging method of lithium-ion battery based on multi-objective BBO algorithm," *Journal of Energy Storage*, vol. 91, p. 112046, 2024. doi: 10.1016/j.est.2024.112046.

[83] Y. Si, X. Zhou, Y. Wang, T. Zhang, Y. Xiao, Y. Liu, "Internal temperature sensing for lithium-ion battery under ultrahigh-rate discharge conditions using physics-informed neural network," *Applied Energy*, vol. 408, p. 127394, 2026. doi: 10.1016/j.apenergy.2026.127394.

[84] X. Tang, Q. Guo, M. Li, C. Wei, Z. Pan, Y. Wang, "Performance analysis on liquid-cooled battery thermal management for electric vehicles based on machine learning," *Journal of Power Sources*, vol. 494, p. 229727, 2021. doi: 10.1016/j.jpowsour.2021.229727.

[85] S. Chen, G. Zhang, J. Zhu, X. Feng, X. Wei, M. Ouyang et al., "Multi-objective optimization design and experimental investigation for a parallel liquid cooling-based Lithium-ion battery module under fast charging," *Applied Thermal Engineering*, vol. 211, p. 118503, 2022. doi: 10.1016/j.applthermaleng.2022.118503.

[86] K. Nam, C. Ahn, "Energy-Efficient Battery Thermal Management in Electric Vehicles Using Artificial-Neural-Network-Based Model Predictive Control," *World Electric Vehicle Journal*, vol. 16, no. 5, p. 279, 2025. doi: 10.3390/wevj16050279.

[87] C. Zhang, S. Zhao, Z. Yang, Y. He, "A multi-fault diagnosis method for lithium-ion battery pack using curvilinear Manhattan distance evaluation and voltage difference analysis," *Journal of Energy Storage*, vol. 67, p. 107575, 2023. doi: 10.1016/j.est.2023.107575.

[88] Z. Jia, Z. Wang, Z. Sun, X. Sun, P. Liu, F. Ruzzenenti, "A multi-scenario data-driven approach for anomaly detection in electric vehicle battery systems," *eTransportation*, vol. 24, p. 100418, 2025. doi: 10.1016/j.etran.2025.100418.

[89] Y. Huang, H. Cai, X. Lai, Y. Zheng, X. Han, D. Ren et al., "Deep learning-driven detection of lithium-plating-type defects for battery manufacturing via formation and capacity grading data," *Journal of Energy Chemistry*, vol. 108, p. 536-549, 2025. doi: 10.1016/j.jechem.2025.04.025.

[90] J. Jeong, E. Kwak, J. h. Kim, K. Y. Oh, "Prediction of thermal runaway for a lithium-ion battery through multiphysics-informed DeepONet with virtual data," *eTransportation*, vol. 21, p. 100337, 2024. doi: 10.1016/j.etran.2024.100337.

[91] L. He, A. Tian, T. Ding, K. Dong, Y. Wang, Y. Gao et al., "Lithium-ion battery state estimation based on adaptive physics-informed neural network of electrochemical model," *Measurement*, vol. 257, p. 118985, 2026. doi: 10.1016/j.measurement.2025.118985.

[92] Z. Xu, Y. Guo, J. H. Saleh, "A physics-informed dynamic deep autoencoder for accurate state-of-health prediction of lithium-ion battery," *Neural Computing and Applications*, vol. 34, no. 18, p. 15997-16017, 2022. doi: 10.1007/s00521-022-07291-5.

[93] J. Tao, S. Wang, W. Cao, C. Fernandez, F. Blaabjerg, "A Comprehensive Review of Multiple Physical and Data-Driven Model Fusion Methods for Accurate Lithium-Ion Battery Inner State Factor Estimation," *Batteries*, vol. 10, no. 12, p. 442, 2024. doi: 10.3390/batteries10120442.

[94] F. Wang, Z. Zhao, Z. Zhai, Z. Shang, R. Yan, X. Chen, "Explainability-driven model improvement for SOH estimation of lithium-ion battery," *Reliability Engineering & System Safety*, vol. 232, p. 109046, 2023. doi: 10.1016/j.ress.2022.109046.

[95] K. Liu, M. F. Niri, G. Apachitei, M. Lain, D. Greenwood, J. Marco, "Interpretable machine learning for battery capacities prediction and coating parameters analysis," *Control Engineering Practice*, vol. 124, p. 105202, 2022. doi: 10.1016/j.conengprac.2022.105202.

[96] X. Liu, H. Peng, B. Li, X. Chen, Z. Li, J. Huang et al., "Untangling Degradation Chemistries of Lithium-Sulfur Batteries Through Interpretable Hybrid Machine Learning," *Angewandte Chemie International Edition*, vol. 61, no. 48, Art. no. e202214037, 2022. doi: 10.1002/anie.202214037.

[97] N. D. K. M. Eaty, P. Bagade, "Digital twin for electric vehicle battery management with incremental learning," *Expert Systems with Applications*, vol. 229, p. 120444, 2023. doi: 10.1016/j.eswa.2023.120444.

[98] Y. Yi, C. Xia, C. Feng, W. Zhang, C. Fu, L. Qian et al., "Digital twin-long short-term memory (LSTM) neural network based real-time temperature prediction and degradation model analysis for lithium-ion battery," *Journal of Energy Storage*, vol. 64, p. 107203, 2023. doi: 10.1016/j.est.2023.107203.

[99] Y. Zhang, M. Zhao, "Cloud-based in-situ battery life prediction and classification using machine learning," *Energy Storage Materials*, vol. 57, p. 346-359, 2023. doi: 10.1016/j.ensm.2023.02.035.

[100] M. Faraji-Niri, M. Rashid, J. Sansom, M. Sheikh, D. Widanage, J. Marco, "Accelerated state of health estimation of second life lithium-ion batteries via electrochemical impedance spectroscopy tests and machine learning techniques," *Journal of Energy Storage*, vol. 58, p. 106295, 2023. doi: 10.1016/j.est.2022.106295.

[101] J. Chen, Q. Yao, P. Kollmeyer, C. Vidal, M. Naguib, S. Panchal et al., "Battery state-of-power estimation: A machine learning battery model with numerical searching approach," *Journal of Energy Storage*, vol. 175, p. 123023, 2026. doi: 10.1016/j.est.2026.123023.

[102] E. Galiounas, T. G. Tranter, R. E. Owen, J. B. Robinson, P. R. Shearing, D. J. Brett, "Battery state-of-charge estimation using machine learning analysis of ultrasonic signatures," *Energy and AI*, vol. 10, p. 100188, 2022. doi: 10.1016/j.egyai.2022.100188.

[103] N. Nicodemo, R. Di Rienzo, M. Lagnoni, A. Bertei, F. Baronti, "Estimation of lithium-ion battery electrochemical properties from equivalent circuit model parameters using machine learning," *Journal of Energy Storage*, vol. 99, p. 113257, 2024. doi: 10.1016/j.est.2024.113257.

[104] S. Haghi, M. F. V. Hidalgo, M. F. Niri, R. Daub, J. Marco, "Machine Learning in Lithium-Ion Battery Cell Production: A Comprehensive Mapping Study," *Batteries & Supercaps*, vol. 6, no. 7, Art. no. e202300046, 2023. doi: 10.1002/batt.202300046.

[105] R. Sun, J. Chen, C. Piao, "Battery health features extraction and state of health estimation based on real-time online vehicle driving data," *Journal of Power Sources*, vol. 645, p. 236784, 2025. doi: 10.1016/j.jpowsour.2025.236784.